

ScoreKit: Unifying Notation and Real-Time Rendering in Fountain-Coach

Overview of ScoreKit and AudioTalk Vision

ScoreKit is a proposed Swift library that bridges high-quality musical notation engraving with interactive, real-time score rendering. It is conceived as a crucial component of the Fountain-Coach ecosystem to realize the **AudioTalk** vision, where users can “talk about sound” in natural language and see those descriptions translated into musical notation and expressive playback ¹ ². In essence, ScoreKit would wrap two key capabilities: (A) the **LilyPond** engraving engine for professional sheet music output, and (B) a Swift-native **score rendering engine** inspired by Verovio for dynamic, on-the-fly visualization. By integrating with FountainAI, Fountain-Store, MIDI2, and Teatro, ScoreKit enables a seamless loop from semantic musical ideas to notation and sound, making “**speak music, hear it happen**” more than just a tagline ³.

LilyPond Integration for Engraving

LilyPond is treated as the **source of truth** for notation in AudioTalk ⁴, due to its proven ability to produce publication-quality scores. ScoreKit will provide a Swift **wrapper around LilyPond**, allowing the system to generate or update sheet music programmatically. In the existing Teatro prototype, a simple `LilyScore` struct already demonstrates this approach by writing a `.ly` file and invoking the `lilypond` command-line tool to produce PDF output ⁵. ScoreKit would build on this by handling LilyPond sessions more robustly – for example, managing temporary files, asynchronous rendering, and error capture (so FountainAI can be alerted to notation mistakes). The LilyPond integration ensures any musical idea or AI-generated fragment can be **typeset into standard notation** and exported (PDF, SVG, etc.) for printing or analysis ⁶.

- *Environment requirements:* LilyPond must be installed on the host (or accessible via a service) for ScoreKit to call it ⁷. On Apple platforms like iOS where bundling GPL software is problematic, ScoreKit can fall back to native solutions (e.g. using PDFKit to draw notation) ⁸. In all cases, the LilyPond `.ly` source remains the authoritative representation that can be saved, versioned, or shared. Notational edits requested via AudioTalk (e.g. “*add a crescendo in bars 1–4*”) will ultimately reflect as LilyPond changes – either by generating corresponding LilyPond syntax or by updating an internal model that is re-engraved to LilyPond format. This guarantees that the formal score always stays aligned with the musical intent applied ².

Swift Native Real-Time Score Rendering

While LilyPond provides beautiful output, it’s not designed for real-time interactivity. To give immediate visual feedback and cross-platform support, ScoreKit includes a **Swift-native score renderer** re-engineered from Verovio’s approach. Verovio is a lightweight engraving library (LGPL-licensed) that renders music from formats like MEI or MusicXML into SVG rapidly ⁹. A Swift incarnation of this idea would allow **dynamic score updates** within apps or web views without needing to rerun LilyPond for every change. ScoreKit’s renderer could parse an internal score representation (or a format like MEI/MusicXML derived from the LilyPond content) and draw notation using CoreGraphics or SwiftUI. This

aligns with prior suggestions to use Apple's PDFKit or native drawing instead of shelling out to LilyPond on iOS ⁸ .

- *Capabilities*: The real-time engine would handle layout of staves, notes, and symbols on the fly. For example, when FountainAI applies a slur or dynamic marking via ScoreKit, the change would immediately reflect on the displayed score view (e.g. a tie curve appears, a crescendo hairpin is drawn) without waiting for a PDF regenerate. This interactive rendering also enables **highlighting and animation** – important for coaching or A/B comparisons. (Notably, the Teatro engine already supports animated SVG outputs and storyboards ¹⁰ ¹¹ ; ScoreKit can extend this to highlight sections being discussed or modified in real time.) If Verovio itself can be wrapped or compiled for Swift, that's an option, but a native implementation grants tighter integration with Fountain-Coach's custom format and future extensions. The Swift renderer ensures that **notation becomes responsive**: it's always in sync with the latest user or AI actions, supporting an interactive "score editor" experience embedded in the Fountain ecosystem.

Fountain-Coach Score Format and Semantic Annotations

To power AudioTalk's **semantic control** of music, ScoreKit would introduce a Fountain-Coach-specific score format or data model that enriches traditional notation with **meaningful tags**. This could be realized as an extension of LilyPond or as a separate JSON/Swift struct representation that the system uses internally. For example, LilyPond comments like `%!AudioTalk: warm_timbre` are proposed as a way to embed *timbral intent* in the sheet music ⁴ . ScoreKit can parse and produce these annotations systematically. Each musical element (notes, phrases, dynamics) in the ScoreKit model can carry metadata linking it to semantic descriptors or AudioTalk *macros*. This fulfills the "semantic memory" idea: the system remembers not just the notes, but the **context and intent behind them**.

- *Custom format*: The format might include unique IDs for measures and phrases, allowing FountainAI to reference them unambiguously (e.g. "increase brightness on phrase #ID42"). It should capture structural data (phrases, sections, instrument assignments) alongside semantic labels (like "*shimmer pad effect*" applied to a section). By having a consistent internal representation, ScoreKit can output standard LilyPond for engraving, **and** serve data to other components (e.g. to the AI or UI) without losing the semantic markers. This also makes it possible to store and query the score in Fountain-Store (the memory core) by those tags or IDs. For instance, the system could retrieve all past scores or passages tagged with "legato crescendo" if needed, using Fountain-Store's fast lookup by keywords ¹² . In short, ScoreKit's format bridges the gap between symbolic music data and the "knowledge" of what that music means in coaching or creative terms.

MIDI 2.0 Integration and Audio Playback

A core purpose of ScoreKit is to connect the notated score with actual sound, leveraging **MIDI 2.0 (UMP)** for high-resolution, expressive playback. When ScoreKit processes a score, it can generate a corresponding stream of MIDI 2.0 events (notes and control changes) that reflect the notation exactly – this is crucial for AudioTalk's "notation ↔ performance bridge" ² . Rather than relying on LilyPond's own MIDI output (which is limited to MIDI 1.0), ScoreKit will translate notation into **UMP messages** in real time. Every marking on the score can influence these messages: for example, a written crescendo becomes a gradual increase in per-note velocity or volume controllers, and a "**brighter**" annotation might translate into MIDI 2.0 **Per-Note Expression (PNX)** data that raises a filter cutoff for those notes

¹³ . The rich resolution and per-note control of MIDI 2.0 means the playback can closely mirror the nuance of the notation.

- *midi2sampler*: ScoreKit's MIDI output will feed into the **midi2sampler** engine (and other synth backends like Csound or SDLKit) to produce sound ¹⁴ . *midi2sampler* is built to handle per-note controllers, high-definition velocity, and other advanced features of MIDI 2.0, so when ScoreKit sends a note with attached timbre/brilliance data, the sampler will actually alter the sound of that note (e.g. crossfading to a brighter sample layer). Integration here means mapping LilyPond instruments or articulations to specific MIDI presets or profiles – for instance, if the score indicates a violin pizzicato, ScoreKit might ensure the MIDI event selects the correct sample zone for that articulation. By working in tandem with the **Audio engine**, ScoreKit helps ensure that *what you see is what you hear*: the symbolic and acoustic representations stay aligned. This closes the loop envisioned in AudioTalk where a user's spoken request yields *both* an updated score and an updated playback. As an example, when a user says “warm legato crescendo on bars 1–4”, the system (via AudioTalk's intent parsing) produces a plan to add a slur and crescendo marking in LilyPond and to adjust timbre/volume in MIDI. ScoreKit would apply the slur and hairpin to the score, and concurrently emit MIDI 2.0 events for those bars that gradually increase expression and warmth, driving the sampler to realize the sound ¹³ .

Integration with FountainAI and Memory Systems

ScoreKit is designed to be **orchestrated by FountainAI**, the AI agent that mediates user requests and system planning in Fountain-Coach. FountainAI can call into ScoreKit's API to accomplish tasks like composing a new score snippet, inserting a musical phrase, or applying an AudioTalk intent. Because ScoreKit abstracts the low-level details, FountainAI can operate at a higher level – e.g., requesting “add an 8-bar chord progression in G major” and letting ScoreKit generate the LilyPond structure for it. This division of labor lets the AI focus on *what* to create, while ScoreKit handles *how* to implement it in notation and data. Crucially, ScoreKit can also feed information back to the AI: for instance, FountainAI could query ScoreKit's model to find out *what dynamics are present* or *how long a section is*, aiding the AI's reasoning or conversational abilities about the music.

- **Fountain-Store synergy**: All musical changes and creations can be persisted via Fountain-Store, using the identifiers and semantic tags provided by ScoreKit. This means FountainAI gains a kind of long-term memory for music. If a user asks “Recall the motif I hummed yesterday”, FountainAI could search the Store (which contains records of a LilyPond snippet or a ScoreKit object from yesterday, tagged accordingly) and retrieve it for editing or replay. Fountain-Store's design ensures that even as the musical knowledge base grows (potentially thousands of phrases, scores, lessons), retrieval remains fast and flexible ¹⁵ ¹² . ScoreKit contributes to this memory by packaging musical information with metadata, making each entry more searchable by content (notes, chords) *and* by concept (mood, technique, tags). In the long run, this integration imbues the system with a **musical semantic memory** – the AI can remember and reuse musical ideas in a coherent way, strengthening the coaching and creative workflow.

UI Integration and Teatro Evolution

On the user interface side, ScoreKit would be the backbone for musical score display and interaction in Fountain-Coach applications. It naturally extends the Teatro engine, which already laid groundwork for rendering and even animating visual elements (including music) in Swift ¹⁰ ¹¹ . By plugging ScoreKit into a UI layer (such as a SwiftUI view or web canvas), users could see their music **rendered in real time** as they or the AI make changes. This interactive score viewer could support features like: click-to-select a bar (for applying an AudioTalk phrase to that region), hover to see semantic annotations (e.g., a

tooltip showing “ brighter timbre” on a note that had that tag), or scrub playback synchronized to the notation. Teatro’s concept of an `Animator` and storyboard could even be employed to highlight changes – for instance, when an AI adds a crescendo, the UI might briefly highlight those bars to draw the user’s attention.

- **Teatro and beyond:** In the original Teatro design, GPT-generated LilyPond was immediately compiled and previewed ⁶ – ScoreKit takes this to the next level by making the preview live and editable. It effectively transforms the static PDF generation into an interactive canvas. Moreover, ScoreKit’s renderer could output not just to screen but also to SVG or other formats for web embedding, aligning with Teatro’s multi-format rendering approach (HTML, Markdown, SVG, image) ¹⁰. Developers can integrate ScoreKit as a Swift package in macOS/iOS apps or use it on a server to generate notation images on-the-fly for a web client. This versatility means any Fountain-Coach interface (a coaching app, a web editor, an AR music display) can all rely on the **same core logic** for consistent results. In short, ScoreKit helps present a unified experience: whether a user is viewing the score, hearing it, or discussing it with the AI, all those modalities stay in sync through this central notation engine.

Conclusion: Bridging Music, Code, and Language

Implementing ScoreKit would be a decisive step toward making the AudioTalk vision a reality. It provides the connective tissue between **symbolic music representation and semantic interpretation**, enabling the system to not only engrave beautiful scores but also understand and manipulate them in context. With LilyPond integration, musicians and AI get authoritative notation output; with a real-time Swift renderer, they get immediacy and interactivity; with a semantic-rich format tied into FountainAI and midi2sampler, every musical idea can propagate through the system – from plain language description to written score to actual sound – in one cohesive loop. ScoreKit thus turns musical *knowledge* into something actionable and persistent within Fountain-Coach. By combining notation, audio, and AI-driven semantics, the Fountain ecosystem moves closer to a future where **language becomes the API for music** ¹⁶, and where creative intent flows effortlessly from the mind to the score to the ears.

Sources: The design draws on Fountain-Coach’s AudioTalk vision ¹ ², existing Teatro rendering prototypes ⁵ ⁶, and the integration of MIDI2.0 and semantic memory across the platform ¹⁴ ¹⁷. These guide the engineering of ScoreKit as a foundational piece of the Fountain-Coach architecture.

1 2 3 4 13 14 16 **VISION.md**

<https://github.com/Fountain-Coach/AudioTalk/blob/7c6b7c3390af8e6cd79f975c074cf6966cfcb70a/VISION.md>

5 6 7 **06_LilyPondMusicRendering.md**

https://github.com/Fountain-Coach/Teatro/blob/21d080f70b2c238469bbb0133a5f14b20afdd0ab/Docs/Chapters/06_LilyPondMusicRendering.md

8 **Addendum.md**

<https://github.com/Fountain-Coach/Teatro/blob/21d080f70b2c238469bbb0133a5f14b20afdd0ab/Docs/Chapters/Addendum.md>

9 **Verovio | Music notation engraving library for MEI with MusicXML ...**

<https://www.verovio.org/>

10 11 **SystemArchitecture.md**

<https://github.com/Fountain-Coach/Teatro/blob/21d080f70b2c238469bbb0133a5f14b20afdd0ab/Docs/SystemArchitecture.md>

12 15 17 **VISION.md**

<https://github.com/Fountain-Coach/Fountain-Store/blob/50654f4d799ff9b5e7bba0c1df57e03c89c56ef1/docs/LEGACY/VISION.md>